



End-to-end transfer cost: large improvements; Torch remains faster

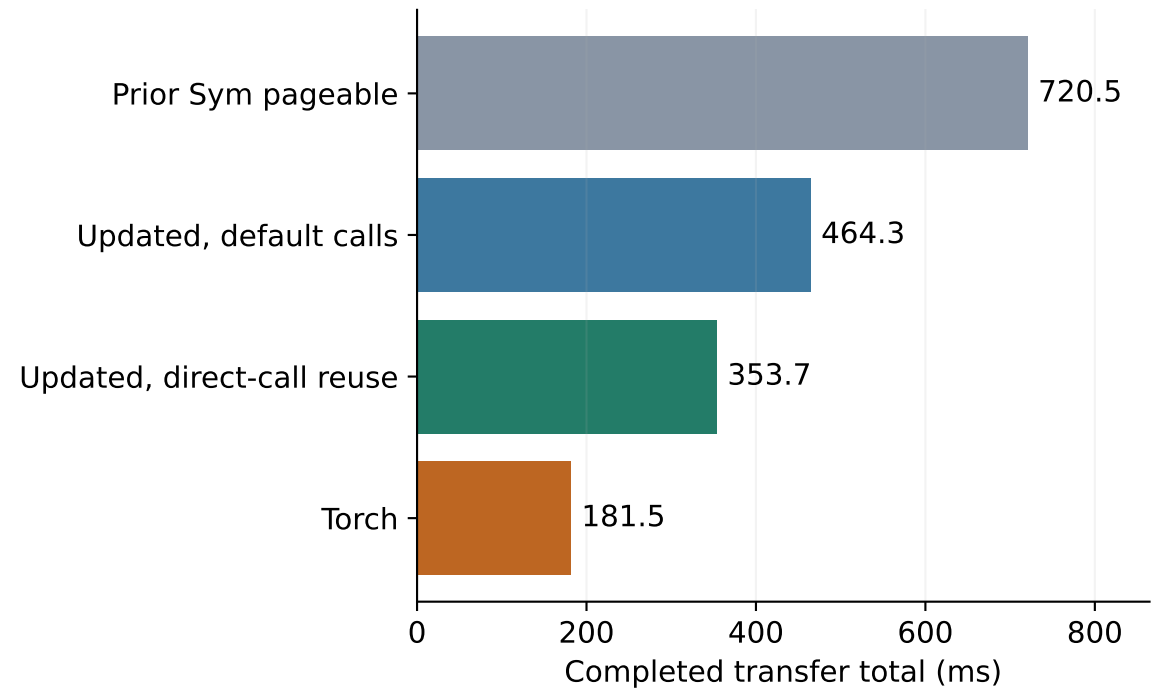
DLRM



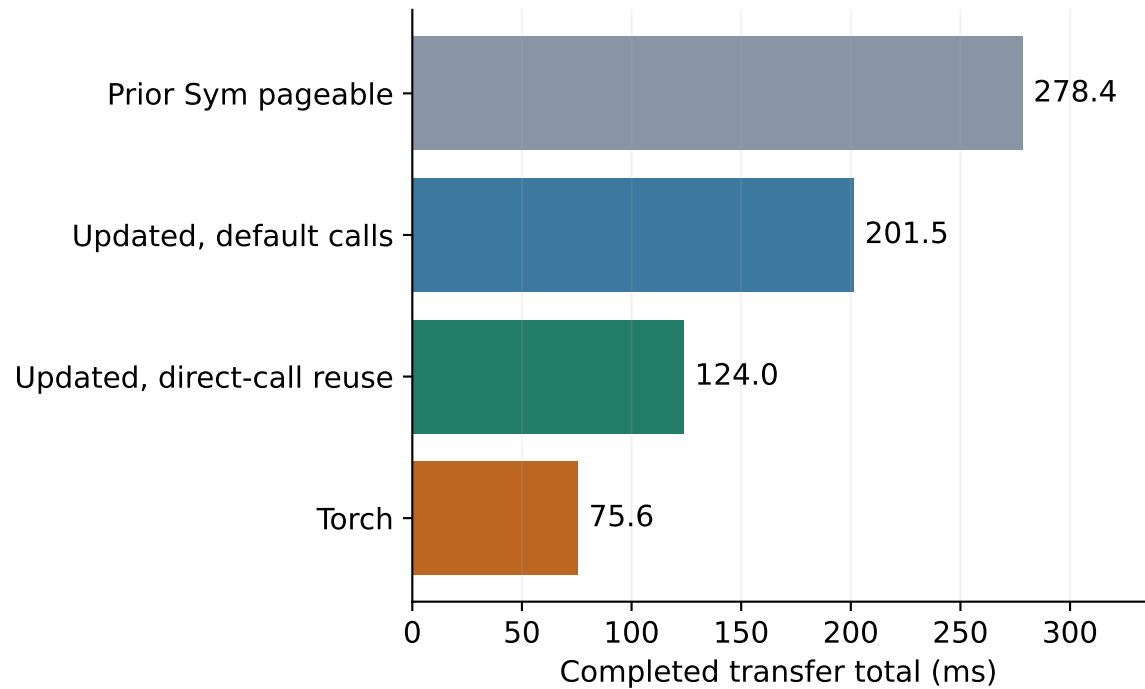
GNN



LLM

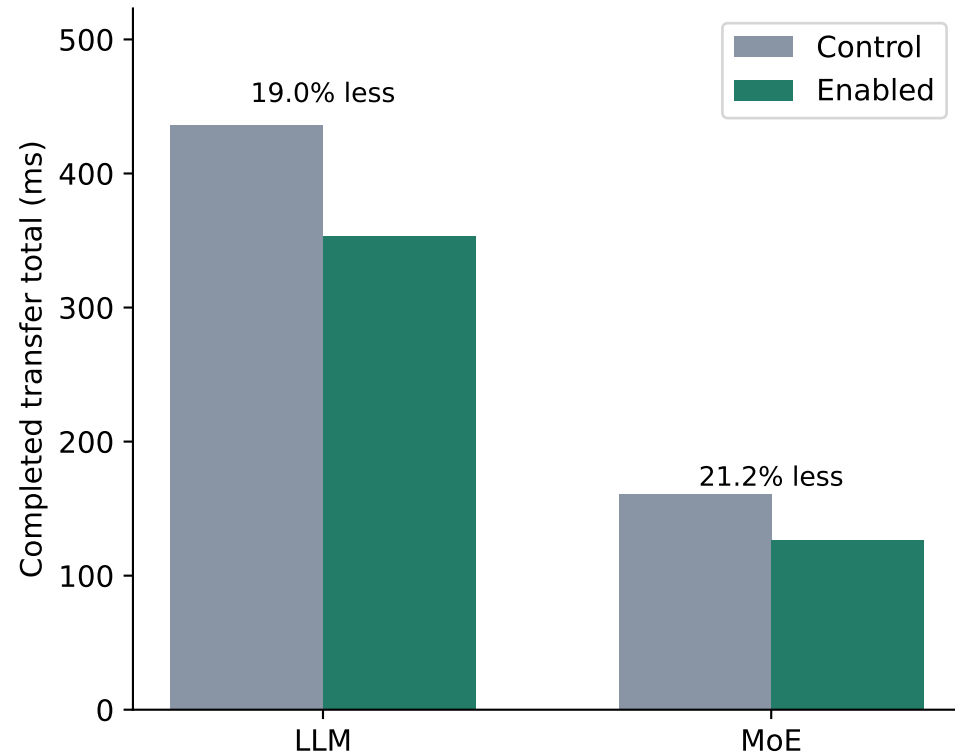


MoE

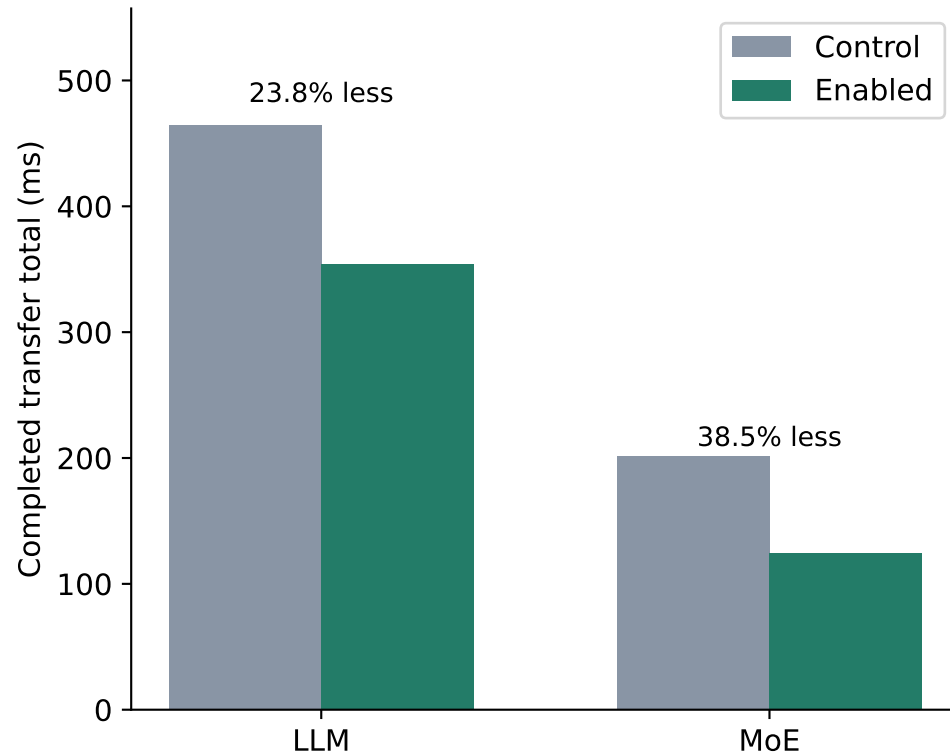


Matched ablations on the updated implementation

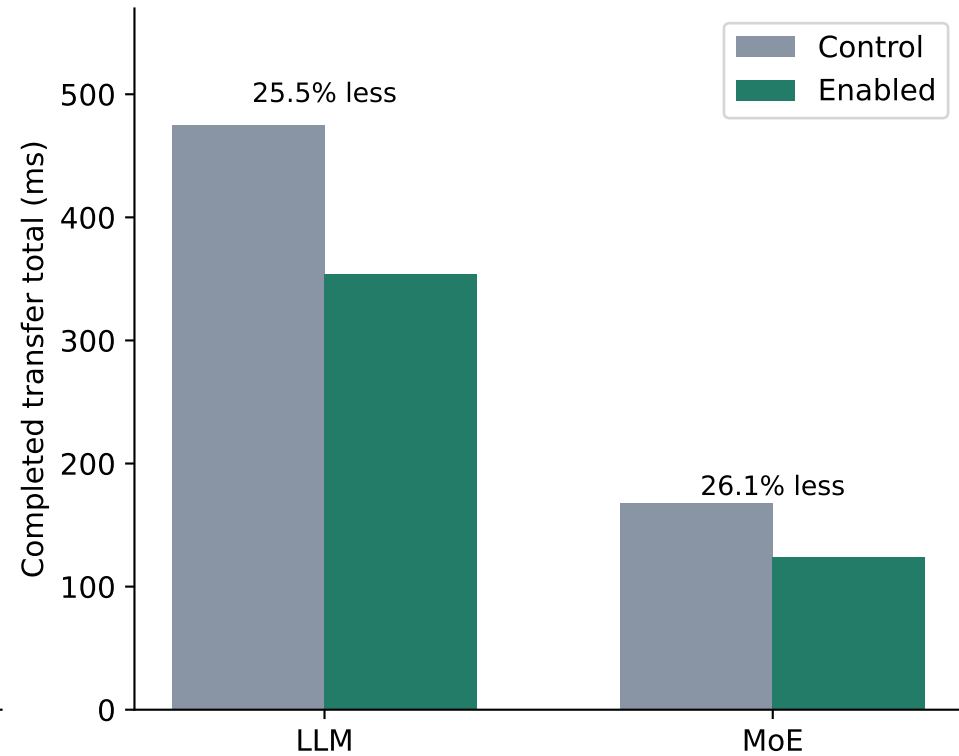
Dense upload: remove Sym staging copy



Direct calls: retain execution resources

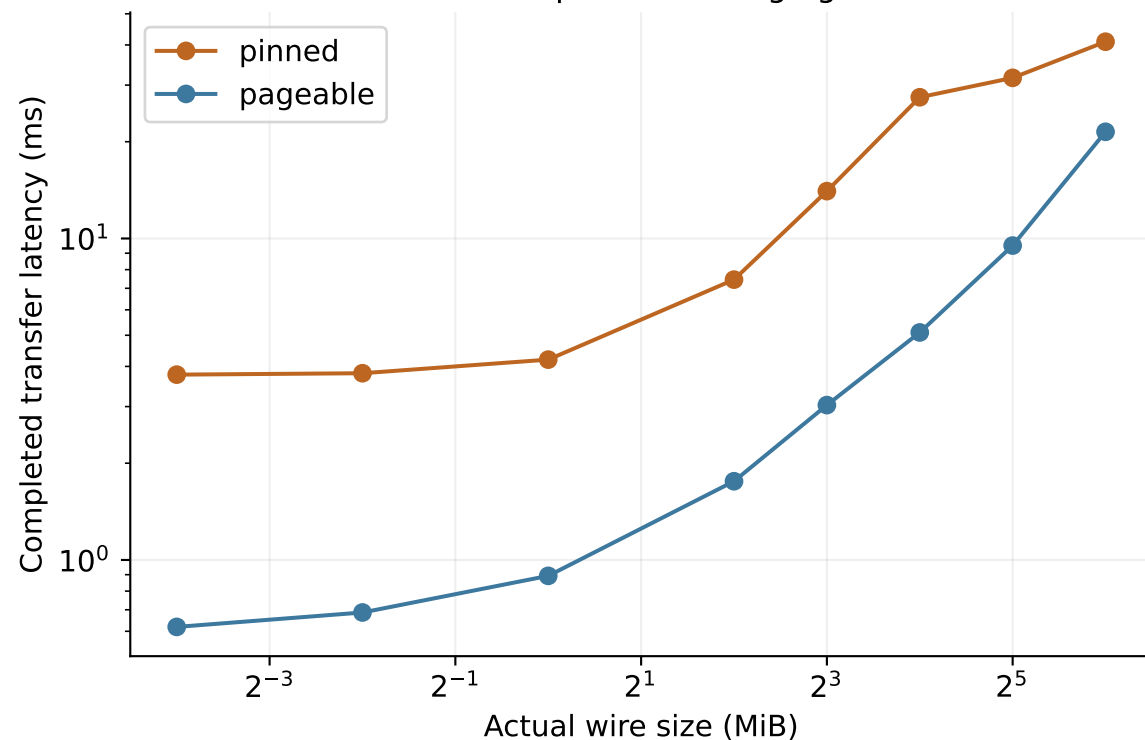


Reuse immutable metadata

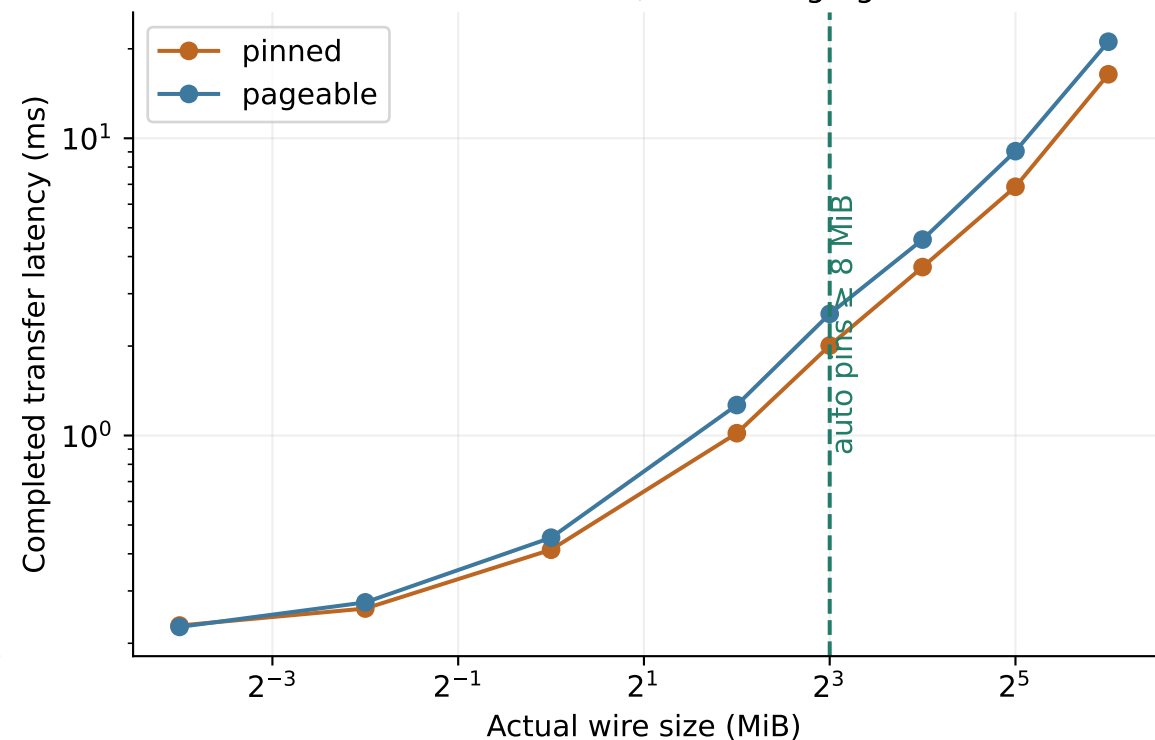


Why auto uses both retention eligibility and size

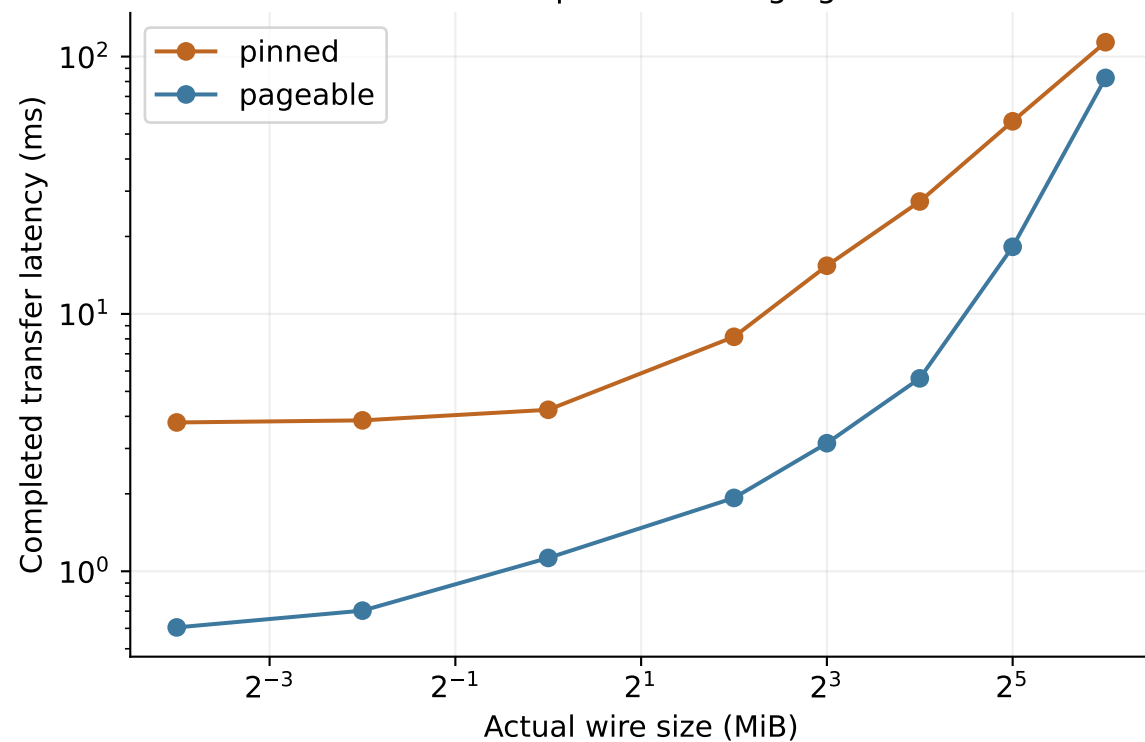
H2D: ephemeral staging



H2D: retained, warm staging



D2H: ephemeral staging



D2H: retained, warm staging

