



Pinning gains depend on reuse: keep first-allocation costs visible



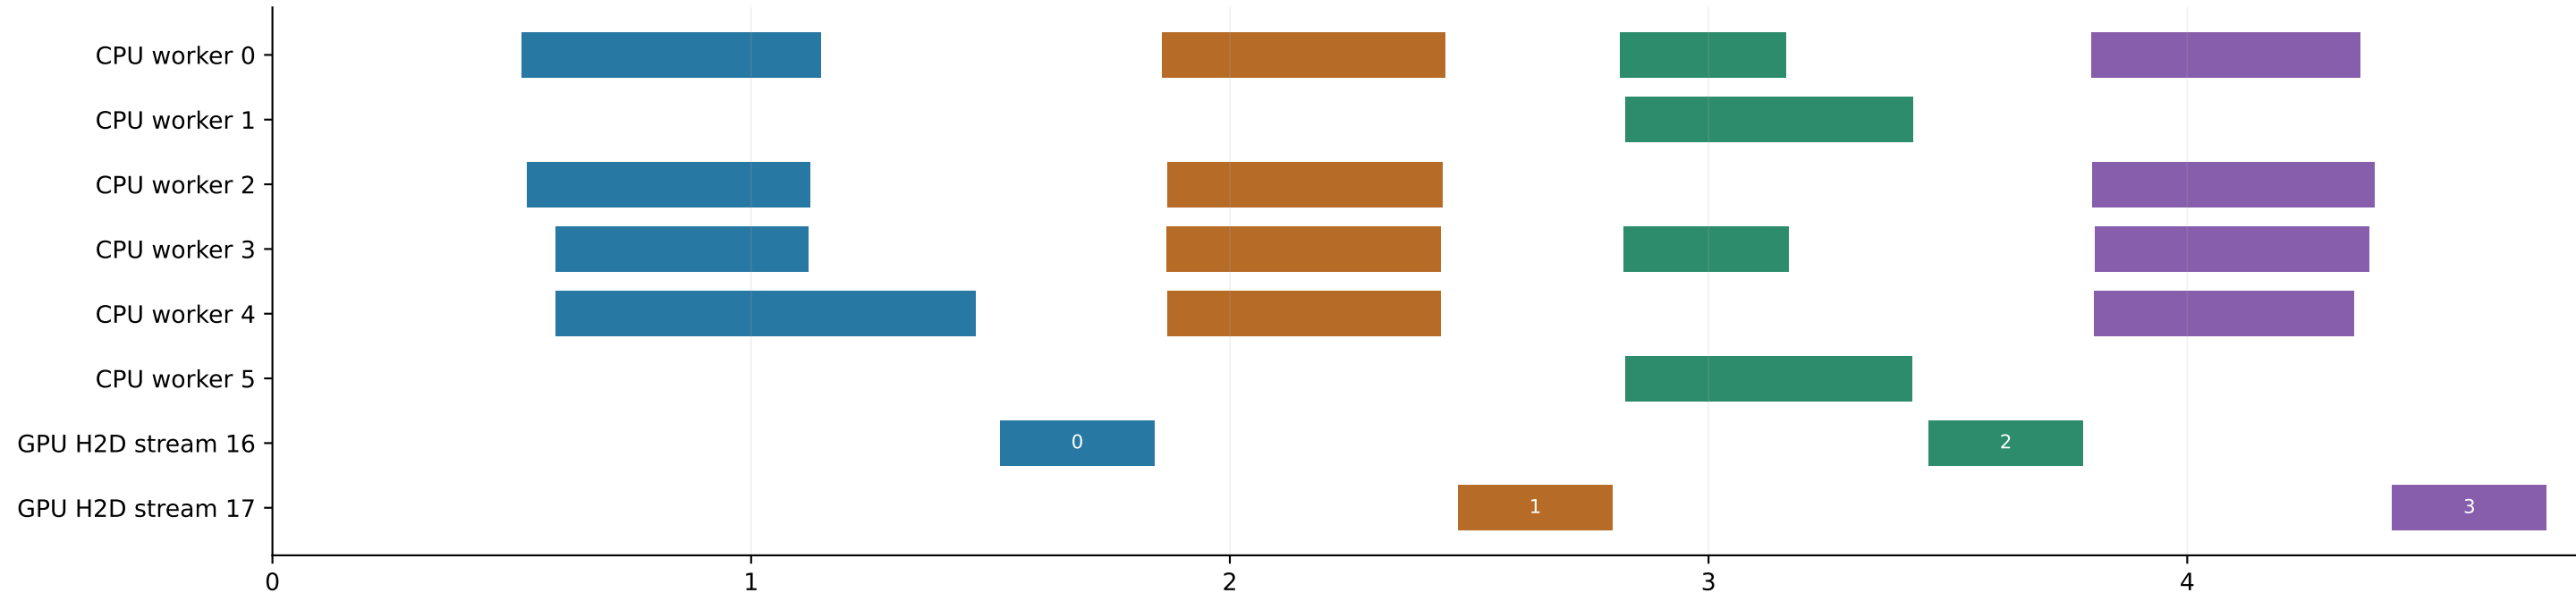
Actual GPU DMA correlated with native CPU gather work (16 MiB, four equal chunks)

configured-4: pinned; gather(n+1)/DMA(n) overlap 0.896 ms



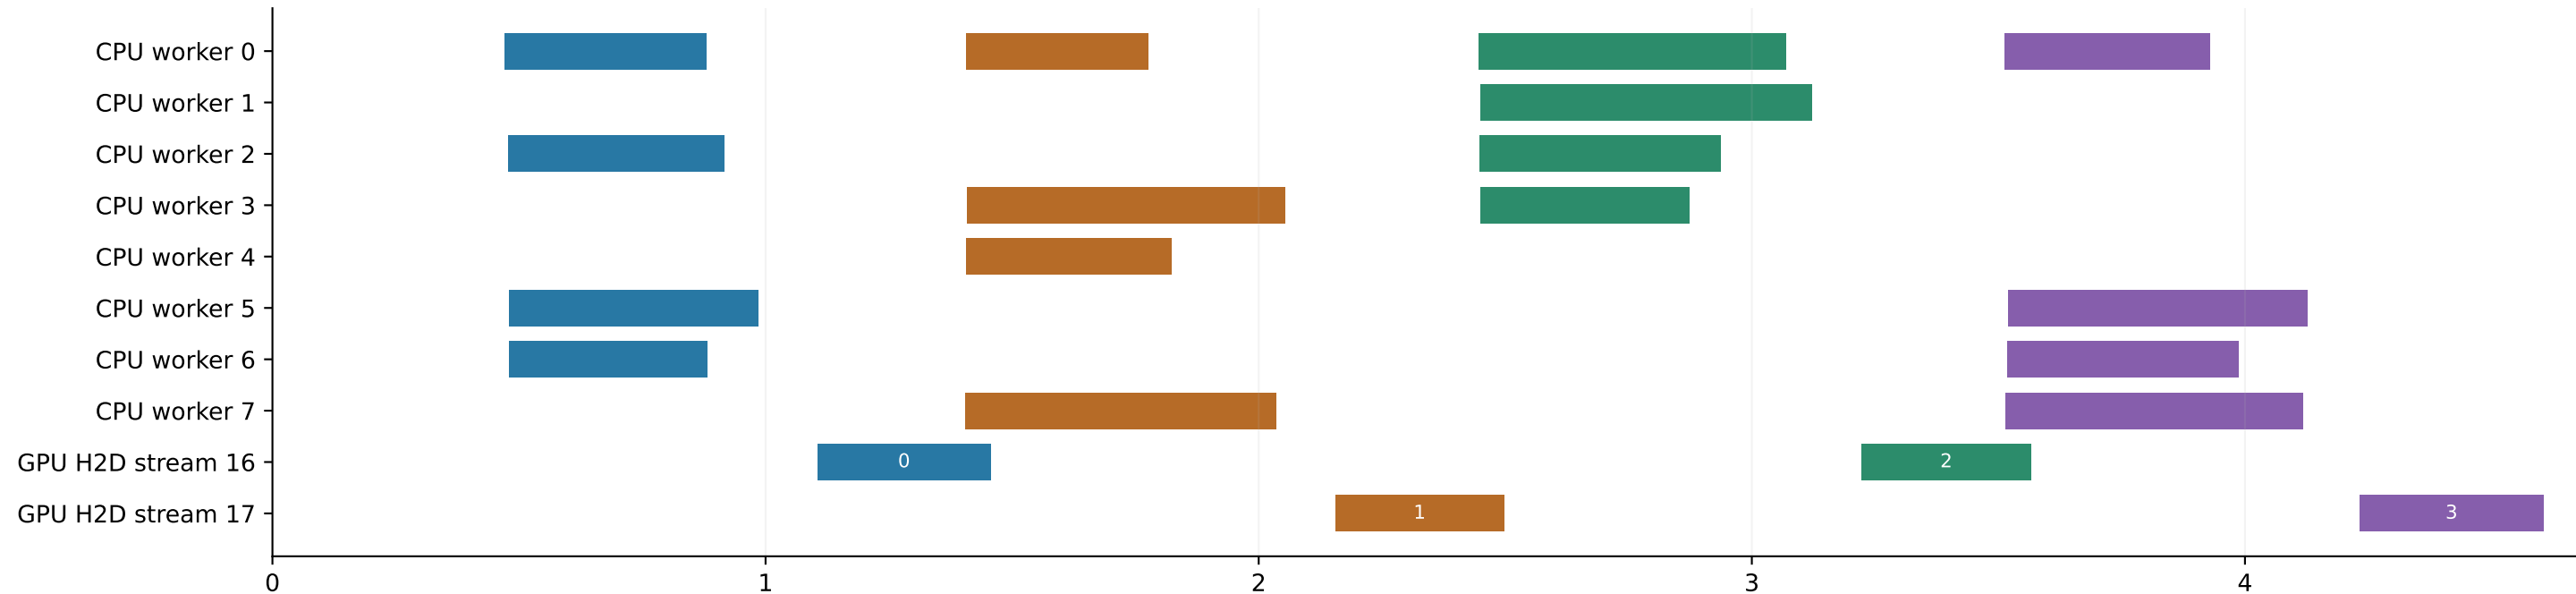
Milliseconds from annotated request start

configured-1: pinned; gather(n+1)/DMA(n) overlap 0.000 ms



Milliseconds from annotated request start

unconfigured-4: pageable; gather(n+1)/DMA(n) overlap 0.156 ms



Milliseconds from annotated request start

Unchanged #162 workloads: policy comparison on the latest implementation

