



LLM: tiny KV calls add up



Sym's GPU conversion is faster



Left: medians per kind, so stacks can differ slightly from median total. Right: separate Nsight runs; kernel time excludes copies.