



The GPU waits for host-side request work

LLM weight fetch · second call, fixed index 1



LLM KV restore · second call, fixed index 1



Illustrative early calls, not median latency; new shapes may need setup. GPU blue = copy; green = kernel. Host colors match attribution.
Separate Torch/Sym calls aligned at entry, not concurrent execution.