



Torch also pays preparation and guard costs


DLRM




GNN



LLM



MOE



Observed preparation / guards CPU transform region CUDA wait API (host)
Output setup / allocation Host/device memory management Other CUDA APIs
Other Python / C++ host work CUDA copy API (host)

Shared classification, disjoint host wall time. Residual host/CPU regions can still contain uninstrumented checks. Diagnostic traces include instrumentation overhead. CPU regions include execution setup; GPU time is not added.