



Inside Torch's native preparation

LLM weight fetch: 5 TensorIterator builds / transfer



KV restore: 2 TensorIterator builds / transfer



Three selected native functions. Additional operand, stride, dimension and output setup work is measured separately.
Sub-microsecond functions are sensitive to observer overhead; the next control exposes that effect.